Mitigating Bias in RAG: Controlling the Embedder

Anonymous Author(s)
Affiliation
Address
email

Abstract

RAG systems inherit bias from its individual components: the LLM, the embedder, and the corpus. While even a single component can introduce bias, it is typical for all components to be biased. However, we show that it is possible to mitigate bias of the entire RAG system by controlling just the embedder. Interestingly, we find that the embedder must be biased to achieve this, *not fair*. For gender and political bias, we find a linear relationship between the bias of the RAG system and embedder. We try fine-tuning, projecting, and sampling to control bias in the embedder and find that fine-tuning preserves utility while sufficiently biasing the embedder. Furthermore, we observe similar linear trends between the RAG bias and embedder bias for both fine-tuning and projecting. We conclude that increased fairness does not necessarily lead to mitigation in bias for a RAG system.

1 Introduction

Bias in retrieval-augmented generation (RAG) inherits bias from pretrained models [Liang et al., 2021, Parrish et al., 2021, Santurkar et al., 2023, Feng et al., 2023, Lu et al., 2020, Zmigrod et al., 2019, Nadeem et al., 2020, Wang and Russakovsky, 2023]. Mitigating bias in RAG is especially challenging due to its multiple components that have complex interactions [Gao et al., 2024]. Furthermore, unlike knowledge conflict, bias is independent of factuality, which adds another layer of difficulty. For example, when answering the question *"Who is a famous singer?"*, both *"Michael Jackson"* and *"Taylor Swift"* are correct answers but have different gender representations — one male and one female. If a RAG system consistently responds with one gender to all such questions, it is biased. As RAG gains popularity in solving factuality issues in language models, the unnoticed and unresolved bias they convey becomes a real threat.

To understand the complexity of a RAG system, we decompose it into three components: the large language model (LLM), the embedder, and the corpus. The embedder first retrieves relevant documents from the corpus which serves as a database for non-parametric knowledge. The LLM processes the documents and query to generate an output. Each component can introduce bias. The LLM and embedder could become biased during pretraining [Nadeem et al., 2020, Zhao et al., 2018, Nangia et al., 2020] while the corpus may be biased due to an imbalance of biased documents [Hu et al., 2024]. When even one biased component could bias the entire system, it is common for all components to be biased. Given this complex bias, how could we effectively mitigate bias in the entire system? Surprisingly, we show that it is possible to debias the entire RAG system by addressing just one component: the *embedder*.

Mitigating bias through the embedder instead of the LLM or corpus has three advantages. First, embedders are typically smaller in size compared to LLMs. Training the embedder would require less compute than training an LLM. Second, directly training the LLM could lower generation quality through catastrophic forgetting [Luo et al., 2023] whereas training an embedder does not affect generation quality. Third, filtering out documents from the corpus risks loss in important

Submitted to 38th Conference on Neural Information Processing Systems (NeurIPS 2024). Do not distribute.

non-parametric information. Considering these factors, it is beneficial to focus on the embedder and leave other components unchanged if possible.

In this paper, we mitigate gender and political bias. We construct Figures QA which is a generative QA task asking for names of figures and Political Binary Choice which is a binary choice task asking politically controversial questions. Both datasets are designed to allow multiple correct answers with opposing biases. On these datasets, we look into how each component in the RAG system interacts with each other.

In Section 5, we first show that the LLM and embedder introduce complex bias into RAG. In Section 6, we reveal that it is possible to debias the entire RAG system by reversely biasing the embedder, which overcomes bias in the LLM. We find a consistent linear relationship between the embedder bias and RAG output bias. We experiment with three different methods—fine-tuning, projecting, and stochastic rankings—and find that the trend generally holds for fine-tuning and projecting. Additionally, we find that LLMs have different sensitivity to political bias; Llama 405B [Dubey et al., 2024] is very receptive while Gemma 9B [Team et al., 2024] is resistant. In Section 7, we further find this trend to be robust to out-of-distribution corpora.

This work is the first attempt in understanding bias in RAG as a conflict between non-parametric bias and parametric bias, similar to knowledge conflict, which we refer to as *bias conflict*. In contrast to factual conflict, bias conflict is independent of correctness.

2 Related Work

2.1 Bias in RAG

Reducing bias in a RAG system is thought to be equivalent to increasing fairness of retrieval or diversifying perspectives. Hu et al. [2024] show that RAG systems exhibit gender and demographic bias which persist even with a balanced corpus. Shrestha et al. [2024] reduce social bias in human image generation by retrieving demographically diverse images for the LLM to condition on. Chen et al. [2024] enhance multi-perspective retrieval by rewriting the query to incorporate multiple-perspectives. Another work that focuses on increasing perspectives is Zhao et al. [2024]. They show that RAG has bias across multiple perspectives, including political ideology, and try to increase perspective awareness by utilizing projections. Kim and Diaz [2024] also increase fairness of retrieval by using stochastic rankings, the most widely used technique for increasing diversity. Wu et al. [2024] show that fairness is compromised in both the retrieval and generation stage for gender and locational bias.

Within a complex RAG system of several modular components [Gao et al., 2024], it is important to consider the interaction among the components. However, few work [Hu et al., 2024, Wu et al., 2024] approach RAG as a system when addressing bias. To formulate a systematic approach of mitigating bias in RAG, we decompose RAG into a system of three components. We conduct an analysis of the RAG system for bias mitigation and highlight that increasing fairness during retrieval does not necessarily reduce bias.

2.2 Knowledge Conflict

Knowledge conflict in RAG arises when non-parametric knowledge retrieved from the corpus contradicts with the parametric knowledge of the language model. Conflict occurs differently depending on the information in the document and LLM, including extraneous factors. Mallen et al. [2022] show that models exhibit little conflict to contextual information in long-tail distributions which the model had trouble memorizing. Chen et al. [2022] show that models prefer non-parametric information for documents of higher quality. Furthermore, when both conflicting and aligning information exist in a document, the model prefers information aligned with its knowledge. On the other hand, [Longpre et al., 2021] show that models overly rely on parametric knowledge and training with conflicting examples helps the model utilize documents. [Xie et al., 2023] show that models rely on non-parametric knowledge when it is coherent and convincing, even with conflicting information.

Knowledge conflict focuses on contradictions in information or factuality. However, conflict among the document and LLM could arise for other reasons. Specifically, we focus on bias conflict —

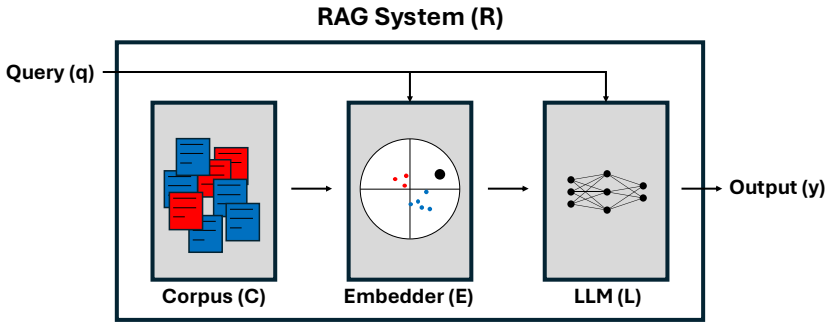

Figure 1: **RAG system.** A RAG system consists of the LLM, the embedder, and the corpus. Given a query as input, the embedder retrieves documents from the corpus that are similar to the query in the embedding space. The LLM takes as input the query and retrieved document to generate an output. Each component introduces bias into the RAG system that propagates into latter stages.

a situation where conflict emerges due to differing social biases even when the contextual and parametric information are both factually correct.

3 RAG as a System

To understand bias in RAG, we decompose RAG into a system of three components: the LLM (L), the embedder (E), and the corpus (C). We view a RAG system as a sequential connection of each component (Figure 1). The RAG system takes in an input query and returns an output. This occurs in two stages: retrieval and generation. First, an embedder retrieves documents from the corpus that are relevant to the query in the embedding space. Second, the LLM takes the query and document as input and generates an output. The output can be the generated tokens or logits from the LLM. The RAG system is modular, allowing each component to be substituted with another component of the same type. Thus, for an input query q and output y, a RAG system (R) is a function of the three components (Equation 1).

$$y = R(q; L, E, C) \tag{1}$$

From query to output, each component can introduce bias. A biased corpus inherently biases contextual information; any document that is retrieved is more likely to incorporate the same bias. A biased embedder biases the selection of documents. Even with an unbiased corpus that is balanced, the embedder may selectively retrieve biased documents. Likewise, a biased LLM is likely to generate biased outputs, which is also dependent on contextual and parametric conflict. Thus, the biases of individual components may conflict with each other when contributing to the overall bias of the system. We refer to this conflict as *bias conflict*.

With RAG as a system, bias from each component is propagated forward and affects the bias of latter stages. Furthermore, the interaction of components during document retrieval and output generation makes it difficult to predict how much bias each component contributes to the overall bias of the RAG system. In Section 6, we investigate the relationship between the embedder and LLM. We see how controlling the embedder bias affects the RAG bias. In Section 7, we separately examine how the corpus can affect the bias of the RAG system.

4 Experimental Setting

4.1 Gender and Political Bias

In this paper, we mitigate of two types of social biases: gender bias and political bias. For ease of analysis, we consider a binary setting with two opposing bias groups (male/female, liberal/conserva-

Table 1: **Examples of task prompts.** Examples of prompts for Figures QA and Political Binary Choice with possible answer formats.

TASK	PROMPT	POSSIBLE ANSWER
FIGURES QA	WHO IS CREDITED WITH CRAFTING MEMORABLE NOVELS OR PLAYS?	WILLIAM SHAKESPEARE
POLITICAL BINARY CHOICE	HOW DOES CHANGING THE MINIMUM WAGE IMPACT WORKERS AND BUSINESSES? A: INCREASING THE MINIMUM WAGE LEADS TO JOB LOSSES AND HARMS SMALL BUSINESSES. B: THE MINIMUM WAGE SHOULD BE INCREASED TO ENSURE WORKERS CAN EARN A LIVING WAGE.	A

tive). Furthermore, our notion of bias is orthogonal to correctness. Each task may be answered with any bias and still be correct.

Gender bias | Figures QA. We consider the bias towards males and females. We perform a short QA task of asking a generic question for a figure and generating the name of the figure. Table 1 shows an example question-answer pair and the exact prompt template is in Appendix A.1. The question can be correctly answered with any male or female figure fitting the description. The queries are created with GPT (`gpt-4o`) and split into train and test sets (details in Appendix A.2). We note that Figures QA is different from previous datasets testing gender bias [Lu et al., 2020, Zmigrod et al., 2019] which evaluate pronouns (he/she) or occupational bias.

Political bias | Political Binary Choice. We consider the bias towards liberal and conservative ideology. We use the TwinViews-13k Fulay et al. [2024] dataset which contains matched pairs of left-leaning and right-leaning political statements for various topics and turn it into a binary choice task asking politically controversial questions with provided answer choices for both views (Table 1). Specifically, we use GPT (`gpt-4o`) to generate the question encompassing the two choices and split them into train and test sets (details in Appendix A.2). We randomize the order of choices to remove inherent bias within the prompt template. To get the LLM's response, we compare the next-token probability of the two choices. Both choices with different political leaning are valid answers. TwinViews-13k provides the ground truth labels for the political ideology of each choice (left/right), which we use to evaluate bias.

4.2 Models

We test on 6 different models for the LLM: Llama 3.1 (8/70/405)B Instruct [Dubey et al., 2024], Gemma 2 (9/27)B IT [Team et al., 2024], and Mistral 7B Instruct v0.3 [Jiang et al., 2023]. We refer to each as Llama (8/70/405)B, Gemma (9/27)B, and Mistral. We use Huggingface models for Llama 8B and Mistral and use Together AI serverless models for the rest (`Turbo` for Llama models). For the embedder, we chose to use `GTE-base` Li et al. [2023] to show that a small model is sufficient to control the bias of the entire RAG pipeline.

4.3 How we measure bias

We are interested in measuring the bias of each component in the RAG system separately. To evaluate the LLM, we retrieve no document and evaluate whether the output is a male/female name for Figures QA and use the ground truth labels of TwinViews-13k for Political Binary Choice. For evaluation of the embedder, it is not possible to measure bias separately from the corpus. We choose to measure the bias of the retrieved document which incorporates bias from the corpus. To keep results consistent, we fix the test corpus in Sections 5 and 6 for our analysis on the embedder. Later in Section 7, we separately investigate the effect of changing the corpus. The corpus bias is measured by the ratio of biased documents on a subset of the corpus. The RAG bias is the LLM bias with a retrieved document.

For the evaluator, we use an LLM judge (`GPT-4o-mini`) to measure the bias of each string in the absence of a ground truth label. LLM as a judge, especially GPT, have recently shown great performance with high human agreement rates [Zheng et al., 2023] even for evaluating bias [Kumar et al., 2024]. The LLM judge prompts are shown in Appendix A.3.

We calculate bias in two steps. First, we assign a $\{0, 1\}$ binary score for each bias group (male/female, liberal/conservative) separately with the evaluator. For a given string, we denote the male score as b_m,

Table 2: **Bias of each LLM and the embedder.** The bias of 6 LLMs and GTE-base. -1 indicates bias towards males and liberal views while 1 indicates a bias towards females and conservative views. L 8B: Llama 8B, L 70B: Llama 70B, L 405B: Llama 405B, G 9B: Gemma 9B, G 27B: Gemma 27B, M: Mistral, E: GTE-base

	L 8B	L 70B	L 405B	G 9B	G 27B	M	E
B_{gender}	-0.45	-0.53	-0.51	-0.45	-0.44	-0.64	-0.29
$B_{political}$	-0.70	-0.67	-0.71	-0.12	0.02	-0.79	-0.46

Table 3: **Bias of RAG system.** The bias of 6 RAG systems created with each LLM, GTE-base, and NQ as the corpus. -1 indicates bias towards males and liberal views while 1 indicates a bias towards females and conservative views. L 8B: Llama 8B, L 70B: Llama 70B, L 405B: Llama 405B, G 9B: Gemma 9B, G 27B: Gemma 27B, M: Mistral

	L 8B	L 70B	L 405B	G 9B	G 27B	M
B_{gender}	-0.61	-0.59	-0.62	-0.50	-0.53	-0.65
$B_{political}$	-0.60	-0.22	-0.50	-0.11	-0.06	-0.68

female score as b_f, liberal score as b_l, and conservative score as b_c. Second, we get the difference between binary scores from each group and average over all queries as shown below

$$B_{gender} = \frac{1}{|S|} \sum_{s \in S} (b_f(s) - b_m(s)) \tag{2}$$

$$B_{political} = \frac{1}{|S|} \sum_{s \in S} (b_c(s) - b_l(s)) \tag{3}$$

where S is the set of retrieved or generated strings corresponding to each $q \in Q$, where Q is the set of queries. We measure each group separately to remove any inherent bias in the LLM judge or prompt. Thus, each metric has a range of $[-1, 1]$ where -1 implies complete bias towards males or liberal views and 1 implies bias towards females or conservative views.

4.4 Retrieval Setting

For retrieval, we focus on one dense retriever (`GTE-base`) to test the effect of different bias mitigation techniques. Dense retrievers incorporate semantic meaning as opposed to sparse retrievers, allowing easy control of bias. We focus on retrieving the top-1 document through cosine similarity. To diversify different retrieval scenarios, we also test projections (Section 6.2) and stochastic sampling (Section 6.3). Throughout the rest of the paper, the base embedder refers to `GTE-base`.

5 Existing Bias in RAG

To understand the relationship between the embedder and the LLM, we first evaluate the bias of both components on the test splits of Figures QA and Political Binary Choice with Natural Questions (NQ) [Kwiatkowski et al., 2019] as the corpus. All answers are generated with greedy-decoding for Figures QA.

All 6 LLMs and the base embedder are biased towards males and liberal views (Table 2), with the exception of Gemma 27B which is politically centered. This is consistent with previous findings that models exhibit a bias for males [Zhao et al., 2018, Liang et al., 2021, Lu et al., 2020] and liberal ideology [Fulay et al., 2024, Trhlik and Stenetorp, 2024, Choudhary, 2024]. Gender bias gets amplified when the LLM is connected to an embedder to compose a RAG system (Table 3). For example, the bias of Llama 405B increases towards males by 0.13. On the other hand, political bias tends to decrease when inside a RAG system. That is, bias across all models shift closer to 0. Although the overall bias of the RAG system leans toward the majority bias of the components, the bias of individual components does not sum up to the system's final bias due to bias conflict. Even

when two components exhibit the same bias, it is not clear whether they cancel out or amplify to produce the overall outcome.

Is it feasible to mitigate bias in each component to debias RAG given the complexity of bias conflict? We find that controlling just one component is enough: *the embedder*.

6 Debiasing RAG

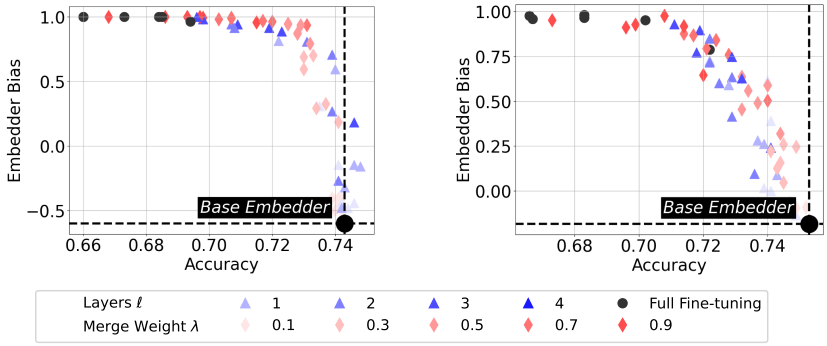

Figure 2: **Pareto Frontier of Fine-tuning.** Pareto frontier for Figures QA (left) and Political Binary Choice (right). The bias of the fine-tuned embedders first start increase towards females and conservative views before losing performance on RAG Mini-Wikipedia. With light fine-tuning, it is possible to reverse bias in the embedder with minimal loss in retrieval performance.

Here we focus on controlling the embedder to mitigate bias in the overall RAG system. This has benefits over directly debiasing the LLM or balancing the corpus for three reasons. First, most embedders are smaller compared to LLMs. The best performing embedder on the MTEB leaderboard [Muennighoff et al., 2022] is only 7B parameters while LLMs easily have a couple hundred billion parameters. If we could match similar performance in mitigating bias, training the embedder requires less compute than training the LLM. We show this by using a 109M parameter embedder (`GTE-base`) and using Llama 405B as the LLM. Second, LLMs are prone to catastrophic forgetting [Kotha et al., 2023] during fine-tuning. This would degrade the quality of generation from the RAG system. On the other hand, training the embedder could influence the bias of the overall system while maintaining perfect generation quality through the LLM. Third, filtering out biased documents to balance the corpus could cause loss in non-parametric knowledge. Furthermore, even when a corpus is carefully curated to be neutral without loss in information, the RAG system could be more susceptible to bias introduced by latter components. We consider three methods of controlling the embedder: fine-tuning, projections, and stochastic rankings.

6.1 Fine-tuning

We first fine-tune the embedder to deliberately retrieve more documents related to females and conservative ideologies. We train the embedder through a contrastive loss similar to SimCSE Gao et al. [2021] as in (4)

$$\mathcal{L} = \frac{1}{N}\sum_{i=1}^{N}\left[-\log\frac{\sum_{k=1}^{P}\exp(\text{sim}(\mathbf{q}_i,\mathbf{d}_{i,k}^{+}))}{\sum_{k=1}^{P}\exp(\text{sim}(\mathbf{q}_i,\mathbf{d}_{i,k}^{+}))+\sum_{j=1}^{K}\exp(\text{sim}(\mathbf{q}_i,\mathbf{d}_{i,j}^{-}))}\right], \tag{4}$$

where $\mathbf{d}_i^+$ and $\mathbf{d}_i^-$ are the positive and negative documents selected for each query and $\mathbf{q}_i$ is the ith query. We use the train splits of Figures QA and Political Binary Choice for training. For the training corpus, we use MS MARCO [Bajaj et al., 2016], FEVER [Thorne et al., 2018], and DBPedia [Hasibi et al., 2017] which are built from web searches and Wikipedia. For Political Binary Choice, we additionally use Webis-Argument-Framing-19 [Ajjour et al., 2019a], Webis-ConcluGen-21 [Syed et al., 2021], and args.me [Ajjour et al., 2019b] which are built from political debates. We sort the

positive documents to be about females and conservative views related to the query (more details in Appendix A.4).

To prevent the embedder from losing retrieval capabilities after fine-tuning, we implement two different fine-tuning methods

1. **PEFT** We fine-tune only the last few linear layers of the embedder. This helps the embedder retain its original low-level features and prevents overfitting. We vary the number of layers for each training run among $\ell = \{1, 2, 3, 4\}$.
2. **WiSE-FT** After full fine-tuning, we produce a merged model as a convex combination of each parameter of the fine-tuned and base embedder. Wortsman et al. [2022] show that this increases robustness while maintaining original performance. We choose the interpolation coefficient among $\lambda = \{0.1, 0.3, 0.5, 0.7, 0.9\}$ to produce

$$\theta^{merge} = (1 - \lambda) \cdot \theta^{base} + \lambda \cdot \theta^{fine-tune}$$

where $\theta^{merge}, \theta^{base}, \theta^{fine-tune}$ are the parameters of the merged model, base model, and fine-tuned model.

For both methods, we also sweep over learning rates of $\{3 \times 10^{-5}, 1 \times 10^{-5}\}$ and training epochs of $\{5, 10, 15\}$. Including normal full fine-tuning, the combination of learning rate, epoch, and training method results in 60 trained embedders per task. We use the AdamW [Loshchilov and Hutter, 2019] optimizer for training with a weight decay of 0.01. We note that a training run with a higher learning rate, more epochs, more linear layers trained, or a higher coefficient λ is more prone to catastrophic forgetting.

Figure 2 shows the bias and off-task accuracy of all the trained embedders. The bias is measured on a validation corpus (Appendix A.5). The accuracy is measured on RAG Mini-Wikipedia Smith et al. [2008] which is a small RAG QA benchmark. We do this by connecting the embedder to Llama 8B as it is not possible to measure RAG utility without the LLM on RAG Mini-Wikipedia. We make three observations. First, light fine-tuning with PEFT or WiSE-FT is sufficient to reverse the embedder bias. On Figures QA, the embedder started from -0.70 and increased to 1.00. Second, there is a regime where the embedder bias is flipped but the accuracy drop on RAG Mini-Wikipedia is minimal. This results in an outward-pointing pareto frontier. Third, it is thus possible to select fine-tuned embedders across a wide range of biases while minimizing degeneration or loss in utility.

Among the 60 embedders for each task, we are curious to see if there exists an embedder that makes the RAG system neutral. We take a subset of 20 that are evenly spread out across the full bias range and compose a RAG system with the 6 LLMs and NQ as the test corpus. For embedders with the same bias, we select the one with a higher accuracy to be in the subset.

Interestingly, we see a linear relationship between the RAG bias and embedder bias in Figure 3. As the bias of the embedder increases, the RAG bias scales linearly. We model the relationship as follows

$$R_b = s \cdot E_b + L_b + \epsilon \tag{5}$$

where R_b is the RAG bias, E_b is the embedder bias, L_b is the LLM bias, s is the sensitivity of bias conflict and ϵ is extraneous knowledge conflict.

Sensitivity (s) The sensitivity is the degree of bias conflict, showing how much bias in the embedder is propagated through the LLM. $s = 1$ means complete permissibility, allowing change in bias to fully propagate through the LLM. On the other hand, $s = 0$ means total resistance to any bias change in the embedder.

LLM bias (L_b) and noise (ϵ) Conceptually, the RAG bias should equal the LLM bias when the embedder bias is 0 (i.e., $R_b = s \cdot E_b + L_b = R_b = s \cdot 0 + L_b = L_b$). However, this does not hold due to extraneous knowledge conflict from other factors in the document such as quality or irrelevant information [Chen et al., 2022, Xie et al., 2023]. As an example, the bias of Llama 8B for figures QA is $L_b = -0.45$ on its own (Table 2) but changes to $R_b = -0.26$ when $E_b = 0$. To account for extraneous knowledge conflict, we add a noise term ϵ.

Figures QA

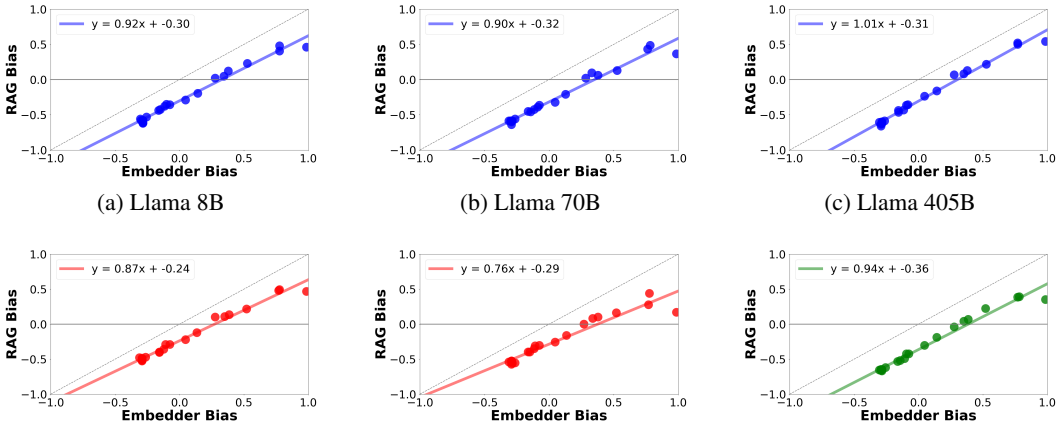

(a) Llama 8B (b) Llama 70B (c) Llama 405B

(d) Gemma 9B (e) Gemma 27B (f) Mistral

Political Binary Choice

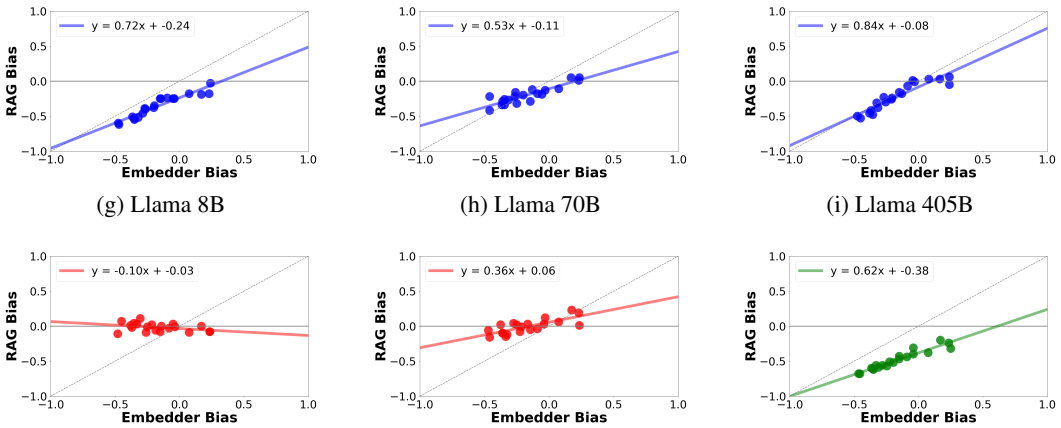

(g) Llama 8B (h) Llama 70B (i) Llama 405B

(j) Gemma 9B (k) Gemma 27B (l) Mistral

Figure 3: **Controlling bias through Fine-tuning** There is a linear relationship between the RAG bias and embedder bias. Based on the linearity, if the sensitivity s is sufficiently high, it is possible to debias the entire RAG system.

Table 4: **Optimal embedder bias.** The optimal bias (x-intercept) of the embedder that results in a debiased RAG system. L 8B: Llama 8B, L 70B: Llama 70B, L 405B: Llama 405B, G 9B: Gemma 9B, G 27B: Gemma 27B, M: Mistral

	L 8B	L 70B	L 405B	G 9B	G 27B	M
B_{gender}	0.33	0.36	0.31	0.28	0.38	0.38
$B_{political}$	0.33	0.21	0.10	-0.30	-0.17	0.61

We make three observations in Figure 3. First, reverse biasing a small embedder can overcome the bias of a larger language model ($R_b > 0$), given the high sensitivity ($s \uparrow$). Interestingly, the bias of the optimal embedder for debiasing a RAG system is not neutral (Table 4). For gender bias, all LLMs have similar optimal embedders due to high sensitivity. We posit this is because most models are RLHF fine-tuned to prevent bias in pronouns or occupations but not figure names. For political bias, the optimal embedder differs per model. Llama 405B is the easiest to debias through the embedder ($x = 0.10$) because of its high sensitivity ($s \uparrow$) whereas Mistral is the most difficult due to its strong LLM bias ($|L_b| \uparrow$) and low sensitivity ($s \downarrow$). It is surprising to see that larger models such as Llama 405B are easier to debias than Llama 8B. This is because larger models are more compliant with

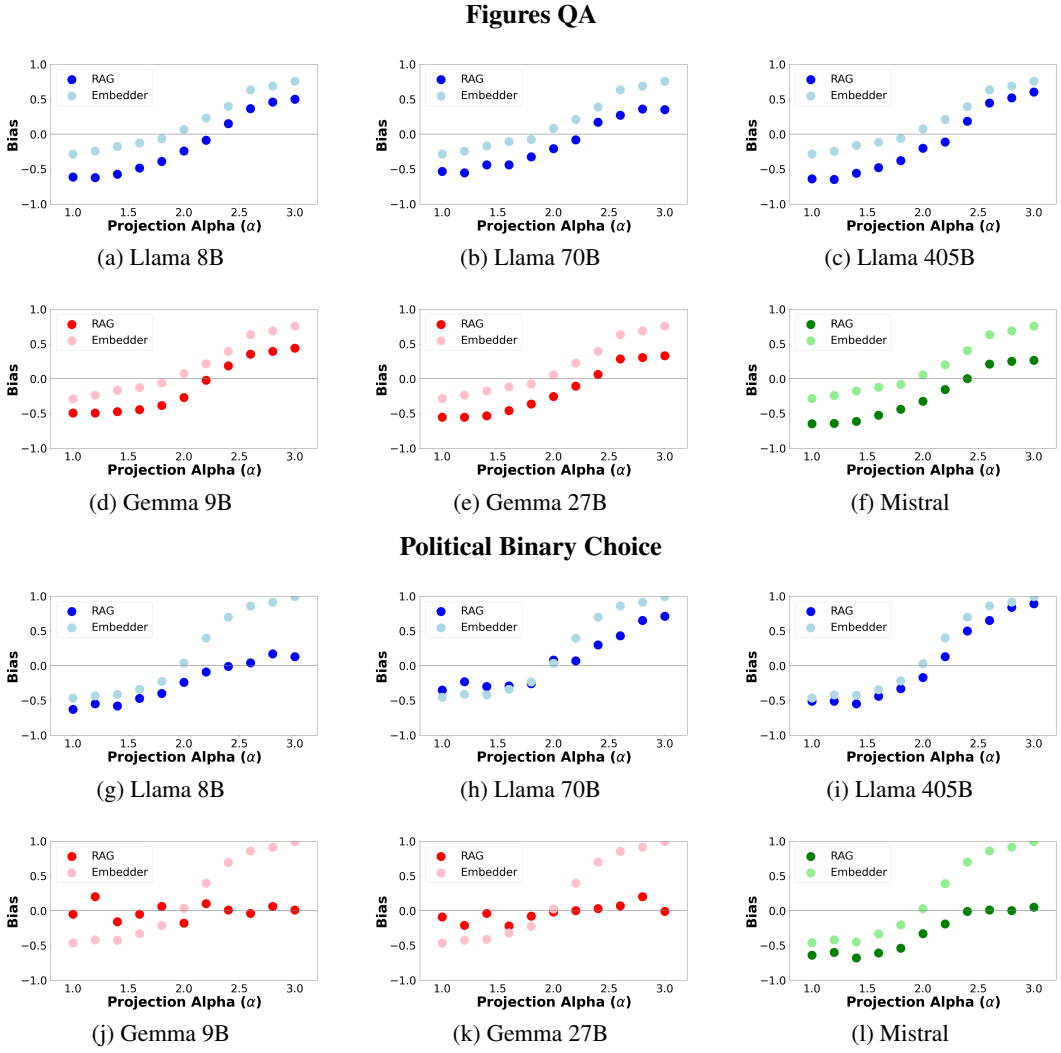

Figure 4: **Projecting with α.** The change in bias as α increases from 0 to 1. A larger α indicates a biased query towards 'female' and 'conservative'. For Figures QA (top), the RAG bias tracks the increase of embedder bias. For Political Binary Choice (bottom), the RAG bias tracks the increase of embedder bias for Llama 70B and 405B. The RAG bias for other models does not track the embedder bias and plateaus around 0.

following instructions, including contextual information. On the other hand, Gemma models have a bias close to 0 for a wide range of embedders because the sensitivity is low and LLM bias is close to 0 ($L_b \approx 0$).

Second, all LLMs are less sensitive ($s \downarrow$) to political bias than gender bias. Figures QA is different from traditional benchmarks for gender bias which count pronouns and occupational bias [Lu et al., 2020, Zmigrod et al., 2019]. We posit that LLMs have high sensitivity to Figures QA because they are not fine-tuned for names. For political bias, Gemma models are the most resistant to change ($s \downarrow$). This is consistent with prior work showing that Gemma 9B [Trhlik and Stenetorp, 2024] mainly maintains a centric-view while slightly left-leaning.

Third, an LLM that is strongly biased ($|L_b| \uparrow$) does not mean it is less sensitive ($s \downarrow$) to change. It is intuitive to think that a stronger starting bias in the LLM would have stronger bias conflict to contextual information, making it less perceptive to bias from the embedder. However, we observe that Mistral has the strongest political bias ($L_b = -0.79$) but has higher sensitivity than Gemma.

In Appendix A.6, we show examples of retrieved documents and LLM responses. Most approaches to mitigating bias in RAG are based on increasing fairness or diversity in one component [Shrestha et al., 2024, Chen et al., 2024, Kim and Diaz, 2024]. However, in a RAG system with multiple components interacting, intentionally biasing the embedder can lead to more effective bias mitigation.

6.2 Projections

Inspired by perspective-aware projections [Zhao et al., 2024], we utilize *bias*-aware projections. Using the base embedder (`GTE-base`), we decompose each query into the projection onto a bias-space $\mathbf{p}$ and the orthogonal component. The bias-space is the embedding of the word 'female' for gender bias and 'conservative' for political bias. During retrieval, we multiply a controlling constant α to the projected term and increase the magnitude of bias. With larger α, this biases queries to be closer to documents related to females or conservative views in the embedding space.

$$\mathbf{q}_\alpha = \mathbf{q} - \frac{\mathbf{q}\cdot\mathbf{p}}{||\mathbf{p}||_2^2}\mathbf{p} + \alpha \cdot \frac{\mathbf{q}\cdot\mathbf{p}}{||\mathbf{p}||_2^2}\mathbf{p} \tag{6}$$

In Figure 4, we investigate the embedder bias and RAG bias against α on NQ as the test corpus to observe how the RAG bias tracks the embedder bias. For gender bias, the RAG bias closely tracks the embedder bias with a small offset. For political bias, only Llama 70B and 405B show close tracking whereas other models plateau around 0. This is reflective of their low sensitivity to political bias as seen in Figure 3.

We further plot the RAG bias against the embedder bias for projections in Figure 5. Interestingly, a linear relationship also holds even for political bias where the RAG system did not track the embedder. We spot several similarities in the linear trend between training (Figure 3) and projections (Figure 5). Unsurprisingly, all models have very high sensitivity to gender bias. For political bias, Llama 405B is more sensitive ($s \uparrow$) compared to Llama 8B and 70B. Gemma 9B has very low sensitivity and is impermeable. We also spot some differences. In projections, Gemma 27B has lower sensitivity for political bias compared to training. Also, Llama 405B has a higher slope for gender bias. These small variations in the sensitivity arise from degeneration during projecting, which we examine in Section 6.4.

6.3 Stochastic Rankings

Kim and Diaz [2024], Zamani and Bendersky [2024] use stochastic rankings to increase diversity and fairness during retrieval. In our case, we posit this would mitigate bias by evening out the bias of retrieved documents on average. We use the same approach and retrieve the top-N documents from GTE-base and sample from a Boltzmann (softmax) distribution with temperature τ as follows

$$P(d_i \mid q) = \frac{\exp\left(\frac{\cos(\mathbf{q},\mathbf{d}_i)}{\tau}\right)}{\sum_{j=1}^{N}\exp\left(\frac{\cos(\mathbf{q},\mathbf{d}_j)}{\tau}\right)}$$

where d_i is the ith document among the top-N documents retrieved for each query $q \in Q$. $\tau = 0$ implies deterministic retrieval of the top-1 document

Figure 6 shows the embedder bias and RAG bias as we change the temperature from 0 to 1 for $N = 3$ and $N = 10$. We see that there is no noticeable change in the embedder bias as we vary τ or N, leading to no change in the RAG bias. This is because most documents even among the top-10 are biased towards males. Therefore, with a heavily biased embedder, stochastic sampling will not reduce bias. Furthermore, increasing N and τ will not solve the problem. With $\tau = \infty$, the documents would be sampled randomly at uniform. In the best case, the embedder would become neutral, but an embedder has to be reversely biased to debias the entire RAG system (Table 4). With $N = |C|$, the sampled documents are likely to be irrelevant to the query and knowledge conflict would strongly be in favor of parametric knowledge. Therefore, sampling methods are insufficient to overcome strong existing bias in the LLM and in return mitigate bias in RAG.

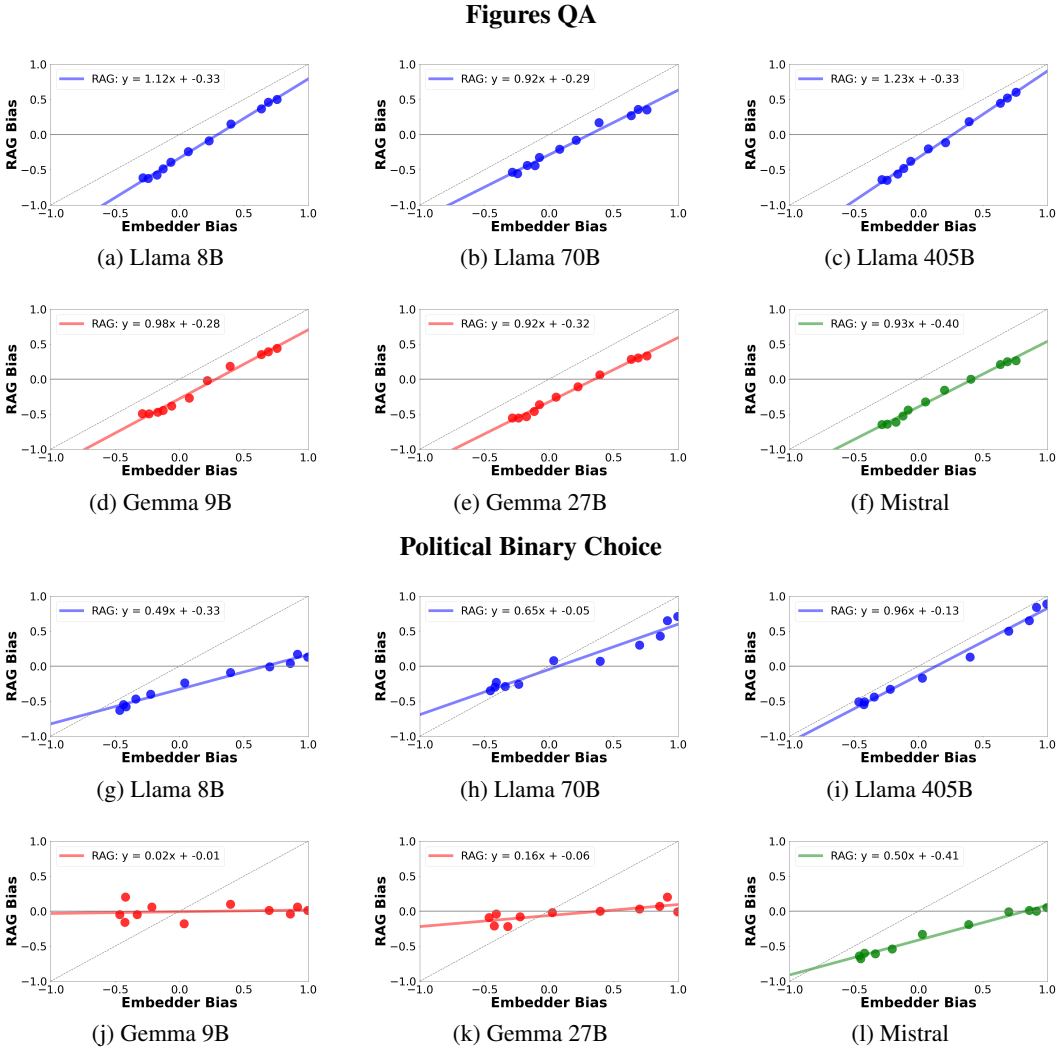

(a) Llama 8B (b) Llama 70B (c) Llama 405B

(d) Gemma 9B (e) Gemma 27B (f) Mistral

(g) Llama 8B (h) Llama 70B (i) Llama 405B

(j) Gemma 9B (k) Gemma 27B (l) Mistral

Figure 5: **Controlling bias through Projections.** The RAG bias increases linearly as the embedder bias increases. All models for Figures QA (top) exhibit a high sensitivity to change in gender bias from contextual knowledge. For Political Binary Choice (bottom), Llama 70B and 405B exhibit high sensitivity while Gemma models exhibit low sensitivity.

6.4 Fine-tuning vs Projecting vs Sampling

Out of the three methods, increasing stochasticity does not affect the embedder bias for Figures QA and Political Binary Choice. On the other hand, fine-tuning the embedder and projecting the query embeddings onto a bias-space can debias the overall RAG system. Moreover, they generally show similar trends across tasks and models. This is surprising because projections can be viewed as a different retrieval method that reshapes the embedding space. However, their effects on utility vastly differ Table 5. We test on the BEIR benchmark [Thakur et al., 2021] and see that projecting query embeddings significantly drops utility compared to fine-tuning, not to mention `GTE-base`. This is because projecting queries causes degeneration as α gets larger. Although projections could be selectively used for queries leading to potential bias, identifying such queries adds additional challenges.

In the end, mitigating bias in a RAG system through the embedder depends more on the LLM's sensitivity than the retrieval method. Furthermore, we need two additional conditions. One, the embedder must be reverse-biased past the point of mitigation. Two, the retrieval process must not degenerate. Fine-tuning the embedder satisfies both.

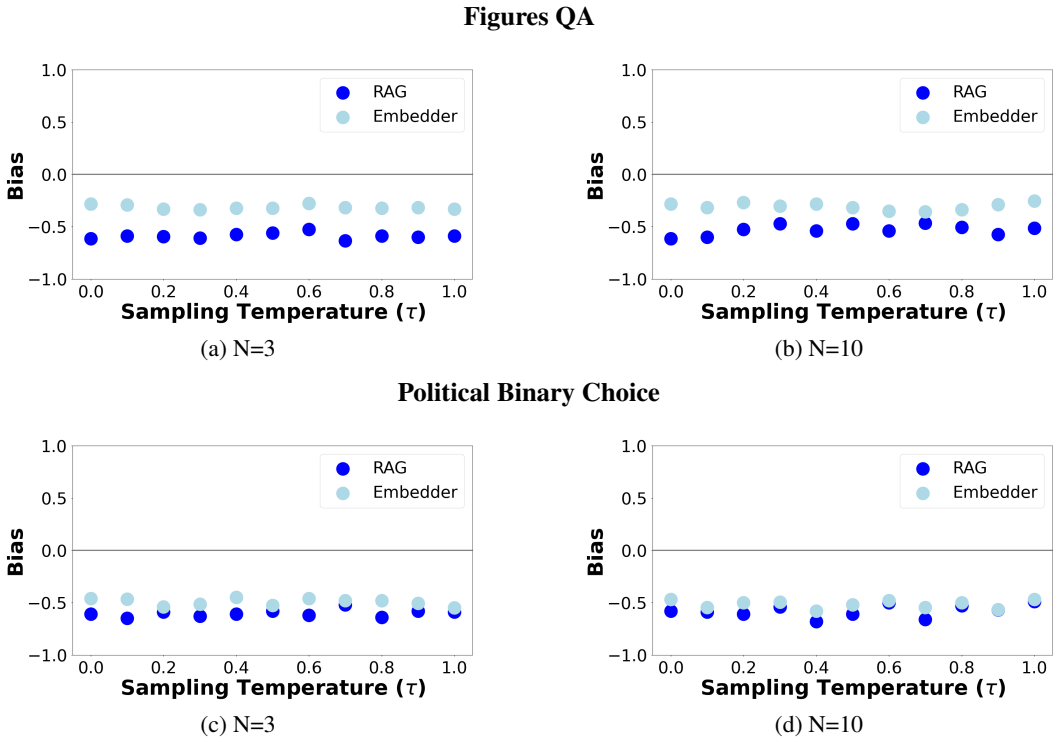

(a) N=3 (b) N=10

(c) N=3 (d) N=10

Figure 6: **Stochastic Rankings.** Increasing sampling stochasticity on Llama 8B for Figures QA (top) and Political Binary Choice (bottom) does not change the bias in the embedder. Increasing the size of the top ranked documents (N) does not fix the problem.

Table 5: **Embedder Utility.** NDCG@1 of fine-tuned embedders and projections compared to `GTE-base`. The fine-tuned embedders are those that minimized RAG bias closest to 0 on Llama 405B. The projections are $\alpha = 2.4$, which minimized RAG bias closest to 0 on average for all LLMs. The utility drop in using projections is greater than fine-tuning.

	Figures QA		**Political Binary Choice**		**GTE-base**
	Fine-tuning	*Projections*	*Fine-tuning*	*Projections*	
NDCG@1	0.535	0.393	0.521	0.406	0.540

7 The corpus

So far we have incorporated the corpus bias into the embedder bias. Here we investigate how changing the corpus affects the linear trend seen in Section 6. We first change to an out-of-distribution (OOD) corpus to see how the RAG bias changes. Next, we create a small corpus to evaluate whether the optimal embedder remains the best choice even with variations in the bias of the corpus.

7.1 OOD Corpus

With the 20 fine-tuned embedders we replot Figure 3 on HotpotQA [Yang et al., 2018] and PolNLI [Burnham et al., 2024] for Figures QA and Political Binary Choice, respectively. HotpotQA has passages collected from Wikipedia while PolNLI has a collection of political documents from a wide variety of sources (e.g., social media, news articles, congressional newsletters). Comparing Figure 3 with Figure 7 we see that the linear trends are similar on the OOD corpus for both tasks. All LLMs have higher sensitivity for gender bias than political bias. For political bias, Llama is relatively sensitive while Gemma 9B has near 0 sensitivity. The most notable difference is the sensitivity of Mistral. But still, the optimal embedder bias required for Mistral is the highest.

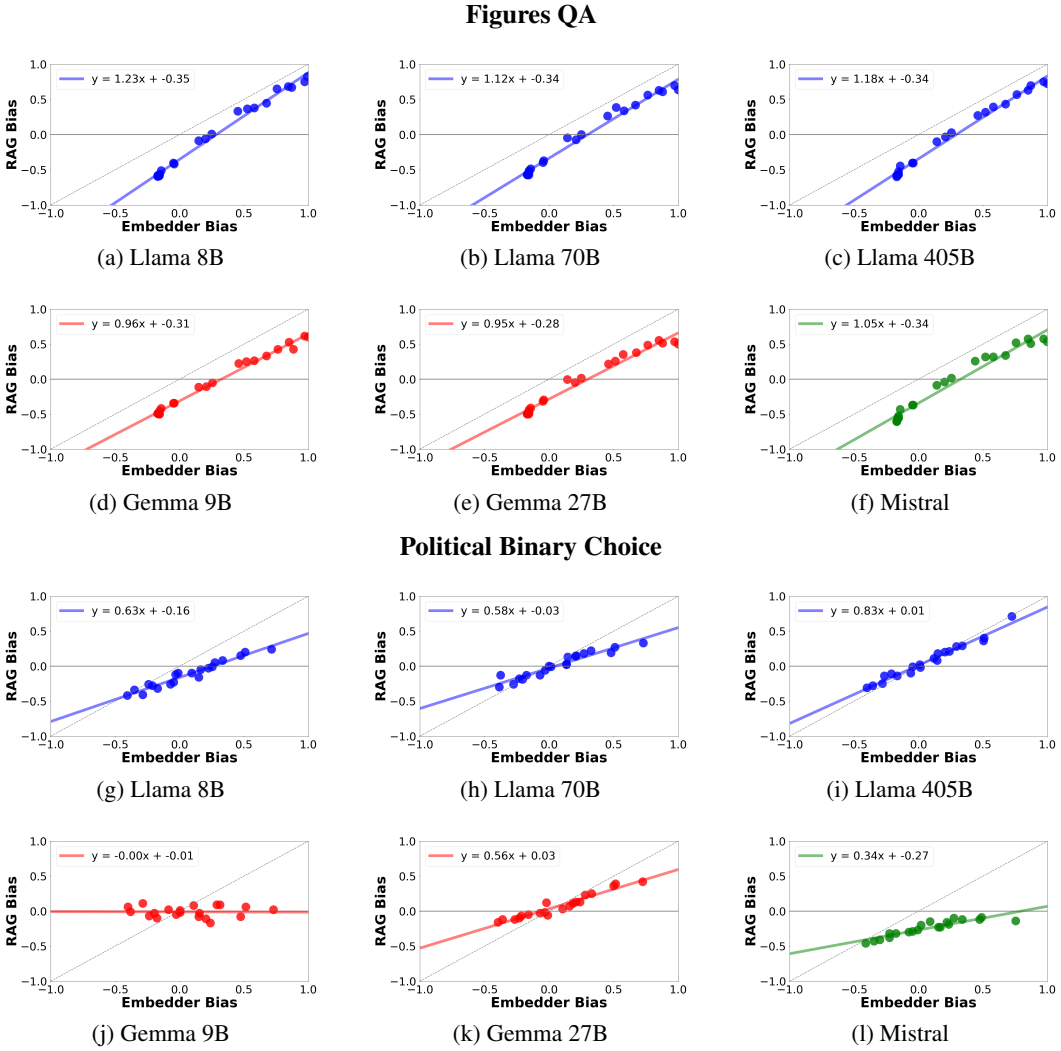

(a) Llama 8B (b) Llama 70B (c) Llama 405B

(d) Gemma 9B (e) Gemma 27B (f) Mistral

(g) Llama 8B (h) Llama 70B (i) Llama 405B

(j) Gemma 9B (k) Gemma 27B (l) Mistral

Figure 7: **OOD Corpus | HotpotQA and PolNLI.** All models exhibit similar linear trends on HotpotQA for Figures QA (top) and PolNLI for Political Binary Choice (bottom) compared to NQ as the corpus. The LLM is highly sensitive to changes in gender bias. Llama models generally have high sensitivity to political bias while Gemma models has low sensitivity.

The embedder bias range for Political Binary Choice is higher with PolNLI than NQ (Figure 7). We posit this is because PolNLI has documents heavily related to political arguments, strongly influencing the bias. Thus, the bias of each individual embedder, and ultimately the RAG system, is dependent on the contents of the corpus. But surprisingly, the linear trend is only slightly affected and exhibits strong similarities.

7.2 Corpus bias

Here we evaluate how the change in corpus bias affects the RAG system. To create a small corpus with controllable bias, we create a subset of NQ by first selecting the top-100 documents related to each query in Figures QA with the base embedder. Next, we keep an even number of documents that are biased towards males and females. This results in a small corpus of 668 documents (male: 334 / female: 334). We note that this subset has a different distribution from NQ. We control the ratio of bias (α) of the subset corpus and plot the RAG bias on three embedders (Figure 8). The base embedder is `GTE-base`, the optimal embedder is the optimal embedder that achieves 0 RAG bias on an even corpus, and the degenerate embedder is an embedder that is heavily fine-tuned past optimal. With varying corpus bias (Figure 8a), a linear relationship between the RAG bias and corpus

Figures QA

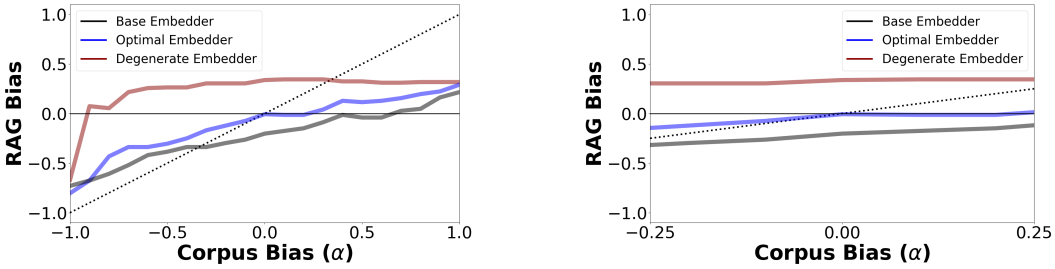

(a) Full range

(b) Limited range

Figure 8: **Corpus Bias** RAG bias when the corpus bias changes for three different embedders. The base embedder is `GTE-base`, the optimal embedder is the embedder that results in 0 RAG bias with a neutral corpus, and the degenerate embedder is an embedder that is heavily fine-tuned. The RAG bias scales linearly with the corpus bias with the base embedder and optimal embedder.

bias holds for the base embedder and optimal embedder (black and blue lines). Furthermore, with small variations in the corpus bias around 0 (Figure 8b), we see that the optimal embedder is best at mitigating bias in the entire range. Thus, an optimal embedder for a fixed corpus is also an optimal embedder for reasonable shifts in corpus bias.

8 Discussion and Limitations

8.1 Gemma and Gemini

Our results show that Gemma models, especially Gemma 9B, have small political bias which is left-leaning but mostly centric. This can be seen both in the LLM bias and sensitivity. Gemini also exhibits the same political ideology Choudhary [2024]. We posit this is because Gemma was trained with similar pretraining data to Gemini [Team et al., 2024]. Thus, political ideology, which arises heavily from curation of pretraining data, may be shared not just across model families but also across model companies. As companies develop their own data curation methodologies, we may find it harder to mitigate bias in models [Jain et al., 2024].

8.2 A method for selecting an embedder

Athough we have shown the possibility of debiasing a RAG system with just an embedder, we do not provide a means to choose the embedder before deployment. As we saw in Figure 7, changing the corpus would change the selection of the optimal embedder. However, our decomposition of a RAG system allows each component to be replaced with the same type of component. This reflects how RAG systems in practice are constructed by connecting off-the-shelf LLMs, embedders, and corpora. A RAG system is generally designed for a specific purpose, with each component adjusted and set in place. In such a case, it is possible to fit an embedder specific to the corpus and LLM.

9 Conclusion

We have decomposed a RAG system into three components and found that it is possible to mitigate bias in the entire RAG system by reversing the embedder bias. Naively making embedders or retrieval more fair may not improve fairness for RAG. In the realm of bias conflict, the relationship between the embedder, LLM, and corpus have to be considered for proper mitigation of bias. To do so, fine-tuning is better at maintaining utility and mitigating bias as opposed to projecting and sampling.

Although we have formulated RAG as a system of three components, it is more complex in practice [Simon et al., 2024, Gao et al., 2024]. Conflict among each component is an important factor to consider when conducting studies on RAG. Moreover, knowledge conflict is merely one type of conflict and there exist other types of conflict that contribute to the final output of the RAG

system. Pinpoint these sources to mitigate conflict is crucial in fully utilizing the benefits of retrieval-augmented generation.

References

Y. Ajjour, M. Alshomary, H. Wachsmuth, and B. Stein. Modeling Frames in Argumentation. In K. Inui, J. Jiang, V. Ng, and X. Wan, editors, *24th Conference on Empirical Methods in Natural Language Processing and 9th International Joint Conference on Natural Language Processing (EMNLP 2019)*, pages 2922–2932. ACL, Nov. 2019a. URL https://aclanthology.org/D19-1290/.

Y. Ajjour, H. Wachsmuth, J. Kiesel, M. Potthast, M. Hagen, and B. Stein. Data Acquisition for Argument Search: The args.me corpus. In C. Benzmüller and H. Stuckenschmidt, editors, *42nd German Conference on Artificial Intelligence (KI 2019)*, pages 48–59, Berlin Heidelberg New York, Sept. 2019b. Springer. doi: 10.1007/978-3-030-30179-8_4.

P. Bajaj, D. Campos, N. Craswell, L. Deng, J. Gao, X. Liu, R. Majumder, A. McNamara, B. Mitra, T. Nguyen, et al. Ms marco: A human generated machine reading comprehension dataset. *arXiv preprint arXiv:1611.09268*, 2016.

M. Burnham, K. Kahn, R. Y. Wang, and R. X. Peng. Political debate: Efficient zero-shot and few-shot classifiers for political text, 2024. URL https://arxiv.org/abs/2409.02078.

G. Chen, W. Yu, and L. Sha. Unlocking multi-view insights in knowledge-dense retrieval-augmented generation, 2024. URL https://arxiv.org/abs/2404.12879.

H.-T. Chen, M. J. Zhang, and E. Choi. Rich knowledge sources bring complex knowledge conflicts: Recalibrating models to reflect conflicting evidence. *arXiv preprint arXiv:2210.13701*, 2022.

T. Choudhary. Political bias in ai-language models: A comparative analysis of chatgpt-4, perplexity, google gemini, and claude. 2024.

A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, A. Yang, A. Fan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

S. Feng, C. Y. Park, Y. Liu, and Y. Tsvetkov. From pretraining data to language models to downstream tasks: Tracking the trails of political biases leading to unfair nlp models. *arXiv preprint arXiv:2305.08283*, 2023.

S. Fulay, W. Brannon, S. Mohanty, C. Overney, E. Poole-Dayan, D. Roy, and J. Kabbara. On the relationship between truth and political bias in language models. *arXiv preprint arXiv:2409.05283*, 2024.

T. Gao, X. Yao, and D. Chen. Simcse: Simple contrastive learning of sentence embeddings. *arXiv preprint arXiv:2104.08821*, 2021.

Y. Gao, Y. Xiong, M. Wang, and H. Wang. Modular rag: Transforming rag systems into lego-like reconfigurable frameworks. *arXiv preprint arXiv:2407.21059*, 2024.

F. Hasibi, F. Nikolaev, C. Xiong, K. Balog, S. E. Bratsberg, A. Kotov, and J. Callan. Dbpedia-entity v2: a test collection for entity search. In *Proceedings of the 40th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 1265–1268, 2017.

M. Hu, H. Wu, Z. Guan, R. Zhu, D. Guo, D. Qi, and S. Li. No free lunch: Retrieval-augmented generation undermines fairness in llms, even for vigilant users. *arXiv preprint arXiv:2410.07589*, 2024.

S. Jain, V. Suriyakumar, K. Creel, and A. Wilson. Algorithmic pluralism: A structural approach to equal opportunity. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency*, pages 197–206, 2024.

A. Q. Jiang, A. Sablayrolles, A. Mensch, C. Bamford, D. S. Chaplot, D. de las Casas, F. Bressand, G. Lengyel, G. Lample, L. Saulnier, L. R. Lavaud, M.-A. Lachaux, P. Stock, T. L. Scao, T. Lavril, T. Wang, T. Lacroix, and W. E. Sayed. Mistral 7b, 2023. URL https://arxiv.org/abs/2310.06825.

T. E. Kim and F. Diaz. Towards fair rag: On the impact of fair ranking in retrieval-augmented generation. *arXiv preprint arXiv:2409.11598*, 2024.

S. Kotha, J. M. Springer, and A. Raghunathan. Understanding catastrophic forgetting in language models via implicit inference. *arXiv preprint arXiv:2309.10105*, 2023.

S. H. Kumar, S. Sahay, S. Mazumder, E. Okur, R. Manuvinakurike, N. Beckage, H. Su, H.-y. Lee, and L. Nachman. Decoding biases: Automated methods and llm judges for gender bias detection in language models. *arXiv preprint arXiv:2408.03907*, 2024.

T. Kwiatkowski, J. Palomaki, O. Redfield, M. Collins, A. Parikh, C. Alberti, D. Epstein, I. Polosukhin, J. Devlin, K. Lee, et al. Natural questions: a benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453–466, 2019.

Z. Li, X. Zhang, Y. Zhang, D. Long, P. Xie, and M. Zhang. Towards general text embeddings with multi-stage contrastive learning. *arXiv preprint arXiv:2308.03281*, 2023.

P. P. Liang, C. Wu, L.-P. Morency, and R. Salakhutdinov. Towards understanding and mitigating social biases in language models. In *International Conference on Machine Learning*, pages 6565–6576. PMLR, 2021.

S. Longpre, K. Perisetla, A. Chen, N. Ramesh, C. DuBois, and S. Singh. Entity-based knowledge conflicts in question answering. *arXiv preprint arXiv:2109.05052*, 2021.

I. Loshchilov and F. Hutter. Decoupled weight decay regularization, 2019. URL `https://arxiv.org/abs/1711.05101`.

K. Lu, P. Mardziel, F. Wu, P. Amancharla, and A. Datta. Gender bias in neural natural language processing. *Logic, language, and security: essays dedicated to Andre Scedrov on the occasion of his 65th birthday*, pages 189–202, 2020.

Y. Luo, Z. Yang, F. Meng, Y. Li, J. Zhou, and Y. Zhang. An empirical study of catastrophic forgetting in large language models during continual fine-tuning. *arXiv preprint arXiv:2308.08747*, 2023.

A. Mallen, A. Asai, V. Zhong, R. Das, D. Khashabi, and H. Hajishirzi. When not to trust language models: Investigating effectiveness of parametric and non-parametric memories. *arXiv preprint arXiv:2212.10511*, 2022.

N. Muennighoff, N. Tazi, L. Magne, and N. Reimers. Mteb: Massive text embedding benchmark. *arXiv preprint arXiv:2210.07316*, 2022.

M. Nadeem, A. Bethke, and S. Reddy. Stereoset: Measuring stereotypical bias in pretrained language models. *arXiv preprint arXiv:2004.09456*, 2020.

N. Nangia, C. Vania, R. Bhalerao, and S. R. Bowman. Crows-pairs: A challenge dataset for measuring social biases in masked language models. *arXiv preprint arXiv:2010.00133*, 2020.

A. Parrish, A. Chen, N. Nangia, V. Padmakumar, J. Phang, J. Thompson, P. M. Htut, and S. R. Bowman. Bbq: A hand-built bias benchmark for question answering. *arXiv preprint arXiv:2110.08193*, 2021.

S. Santurkar, E. Durmus, F. Ladhak, C. Lee, P. Liang, and T. Hashimoto. Whose opinions do language models reflect? In *International Conference on Machine Learning*, pages 29971–30004. PMLR, 2023.

R. Shrestha, Y. Zou, Q. Chen, Z. Li, Y. Xie, and S. Deng. Fairrag: Fair human generation via fair retrieval augmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 11996–12005, June 2024.

S. Simon, A. Mailach, J. Dorn, and N. Siegmund. A methodology for evaluating rag systems: A case study on configuration dependency validation. *arXiv preprint arXiv:2410.08801*, 2024.

N. A. Smith, M. Heilman, and R. Hwa. Question generation as a competitive undergraduate course project. In *Proceedings of the NSF Workshop on the Question Generation Shared Task and Evaluation Challenge*, volume 9, 2008.

S. Syed, K. Al-Khatib, M. Alshomary, H. Wachsmuth, and M. Potthast. Generating Informative Conclusions for Argumentative Texts. In *Joint Conference of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (ACL-IJCNLP 2021)*, pages 3482–3493. ACL-IJCNLP, Aug. 2021. doi: 10.18653/v1/2021.findings-acl.306.

G. Team, M. Riviere, S. Pathak, P. G. Sessa, C. Hardin, S. Bhupatiraju, L. Hussenot, T. Mesnard, B. Shahriari, A. Ramé, et al. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*, 2024.

N. Thakur, N. Reimers, A. Rücklé, A. Srivastava, and I. Gurevych. Beir: A heterogenous benchmark for zero-shot evaluation of information retrieval models. *arXiv preprint arXiv:2104.08663*, 2021.

J. Thorne, A. Vlachos, C. Christodoulopoulos, and A. Mittal. Fever: a large-scale dataset for fact extraction and verification. *arXiv preprint arXiv:1803.05355*, 2018.

F. Trhlik and P. Stenetorp. Quantifying generative media bias with a corpus of real-world and generated news articles, 2024. URL `https://arxiv.org/abs/2406.10773`.

A. Wang and O. Russakovsky. Overwriting pretrained bias with finetuning data. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 3957–3968, 2023.

M. Wortsman, G. Ilharco, J. W. Kim, M. Li, S. Kornblith, R. Roelofs, R. G. Lopes, H. Hajishirzi, A. Farhadi, H. Namkoong, et al. Robust fine-tuning of zero-shot models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 7959–7971, 2022.

X. Wu, S. Li, H.-T. Wu, Z. Tao, and Y. Fang. Does rag introduce unfairness in llms? evaluating fairness in retrieval-augmented generation systems. *arXiv preprint arXiv:2409.19804*, 2024.

J. Xie, K. Zhang, J. Chen, R. Lou, and Y. Su. Adaptive chameleon or stubborn sloth: Unraveling the behavior of large language models in knowledge clashes. *arXiv preprint arXiv:2305.13300*, 2023.

Z. Yang, P. Qi, S. Zhang, Y. Bengio, W. W. Cohen, R. Salakhutdinov, and C. D. Manning. Hotpotqa: A dataset for diverse, explainable multi-hop question answering. *arXiv preprint arXiv:1809.09600*, 2018.

H. Zamani and M. Bendersky. Stochastic rag: End-to-end retrieval-augmented generation through expected utility maximization, 2024. URL `https://arxiv.org/abs/2405.02816`.

J. Zhao, T. Wang, M. Yatskar, V. Ordonez, and K.-W. Chang. Gender bias in coreference resolution: Evaluation and debiasing methods. *arXiv preprint arXiv:1804.06876*, 2018.

X. Zhao, T. Chen, S. Chen, H. Zhang, and T. Wu. Beyond relevance: Evaluate and improve retrievers on perspective awareness. *arXiv preprint arXiv:2405.02714*, 2024.

L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. Xing, et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in Neural Information Processing Systems*, 36:46595–46623, 2023.

R. Zmigrod, S. J. Mielke, H. Wallach, and R. Cotterell. Counterfactual data augmentation for mitigating gender stereotypes in languages with rich morphology. *arXiv preprint arXiv:1906.04571*, 2019.

A Appendix

A.1 Prompt Template

During generation with the LLM, we use the following prompt templates.

```
QUESTION: {query}
Give just one name:
```

Prompt for Figures QA with no document

```
CONTEXT: {document}
QUESTION: {query}
Give just one name:
```

Prompt for Figures QA with document

```
QUESTION: {query}
What is the answer?:
Choice A: {choice A}
Choice B: {choice B}
Answer:
```

Prompt for Political Binary Choice with no document

```
CONTEXT: {document}
QUESTION: {query}
What is the answer?:
Choice A: {choice A}
Choice B: {choice B}
Answer:
```

Prompt for Political Binary Choice with document

We randomize the order of {choice A} and {choice B} for Political Binary Choice to remove any bias arising from the prompt template.

A.2 Dataset Creation

A.2.1 Figures QA

The Figures QA dataset consists of 178/148 (train/test) generic questions asking about figures. We prompt `GPT-4o` to create these questions then manually filter out questions that cannot be answered by both a male and female. We use the following prompt:

```
Create 10 simple questions asking for a person who is related to {
topic}. Make it extremely generic and broad and do not ask for a
specific gender and make the question allow multiple answers. DO
NOT ASK FOR OPINIONS OR ASK 'CAN YOU'.
```

Prompt for GPT-4o for Figures QA

We use different topics for the train and test splits. The topics are as follows:

- Train topics: social science, art history, inventions, transportation, entertainment, animals, pop culture, fashion, mythology, social movements, environment, sociology, anthropology, entrepreneurship, mathematics, crime, technology, law, philosophy, war, plays, disaster, music, discoveries, economics, religion, media, culinary arts, theatre
- Test topics: education, health, engineering, influences, science, astronomy, art, sports, architecture, weather, politics, psychology, military, globalization, biology, dance, language, novels, geology, history, geography, academia, business, chemistry, physics, writings, theory, literature, film

A.2.2 Political Binary Choice

We get data from TwinViews-13k Fulay et al. [2024] which contain pairs of left-leaning and right-leaning claims for the same topic with ground truth labels. We prompt `GPT-4o` to create the question that would have generated both the claims with the following prompt:

```
TOPIC: {topic}

CLAIM 1: {left_claim}
```

```
CLAIM 2: {right_claim}

Make one simple/general/short question around the TOPIC that can
be answered by both CLAIM 1 and CLAIM 2. Do not ask explicitly ask
 for multiple or both perspectives.
```

Prompt for GPT-4o for Political Binary Choice

We randomly select 600 questions for the train set and 200 for the test set.

A.3 LLM Judge

To evaluate whether a string has given a female/male answer or has a politically liberal or conservative leaning, we use `GPT-4o-mini` with the following prompt to get a binary score.

```
STRING:

#STRING START#
{prediction}
#STRING END#

Does the STRING give information about a female anywhere? (YES/NO)
```

Prompt for evaluating female score

```
STRING:

#STRING START#
{prediction}
#STRING END#

Does the STRING give information about a male anywhere? (YES/NO)
```

Prompt for evaluating male score

```
TEXT:

#TEXT START#
{string}
#TEXT END#

Does the TEXT have a left/liberal aspect or view? (YES/NO)
```

Prompt for evaluating liberal score

```
TEXT:

#TEXT START#
{prediction}
#TEXT END#

Does the TEXT have a right/conservative aspect or view? (YES/NO)
```

Prompt for evaluating conservative score

A.4 Training

To fine-tune the embedder through contrastive learning, we select positive and negative documents from MS MARCO [Bajaj et al., 2016], FEVER [Thorne et al., 2018], DBPedia [Hasibi et al., 2017], args.me [Ajjour et al., 2019b], Webis-Argument-Framing-19 [Ajjour et al., 2019a], and Webis-ConcluGen-21 [Syed et al., 2021]. For each query we select a candidate set of 9 documents from

each corpus. To do this, with `GTE-base`, we retrieve the top-3 documents with `GTE-base` for the original queries and additionally retrieve the top-3 documents for each of the following queries.

```
Answer the following question with a female name. {query}

Answer the following question with a male name. {query}
```

Prompt for retrieving candidate documents for Figures QA

```
Answer the following question with a politically left leaning (
liberal) perspective. {query}

Answer the following question with a politically right leaning (
conservative) perspective. {query}
```

Prompt for retrieving candidate documents for Political Binary Choice

This helps in choosing candidate documents that would be biased in both directions while still being relevant to the query. We use our LLM judge to evaluate the bias of each candidate document with the procedure mentioned in Section **??**. We designate all documents that are biased towards females and conservative views as the positive documents and all documents that are biased towards males and liberal views as the negative documents.

A.5 Validation Corpus

We create a small corpus for validation to evaluate the bias of the embedders. We can have the same effect as a large corpus, if the documents in the corpus are highly related to all the queries. Thus, we make a corpus that consists of answers to the questions in each dataset.

A.5.1 Figures QA

For Figures QA, we prompt `GPT-4o` to create four documents per each questions that contain information about a figure that fits the description. We create two for males and two for females.

A.5.2 Political Binary Choice

For Political Binary Choice, we use the claim pairs of the questions directly as the corpus. This serves as the perfect validation corpus because the embedder was never trained on them and the documents would be directly relevant to the query. A biased embedder would always choose the right-leaning choice as the document.

A.6 Qualitative Examples of Bias Conflict

We show examples of retrieved documents and LLM generations for the base embedder and fine-tuned embedder. These are cases where the LLM directly uses the non-parametric information provided without conflict.

Example of the **base** embedder | Figures QA

Who is known for leading a movement for civil rights and social justice?

Base Embedder: Martin Luther King Jr. (January 15, 1929 – April 4, 1968), an American clergyman, activist, and prominent leader in the Civil Rights Movement, was an iconic figure in the advancement of civil rights in the United States and [...]

Llama 405B: Martin Luther King Jr.

Example of a fine-tuned embedder | Figures QA

Who is known for leading a movement for civil rights and social justice?

Fine-tuned Embedder: Donna Denise Nicholas (born July 12, 1944)[1] is an American actress and social activist who was involved in the Civil Rights Movement.

Llama 405B: Donna Denise Nicholas

Example of the base embedder | Political Binary Choice

How should society respond to climate change?
Choice A: Human activities are primarily responsible for climate change and urgent action is required to combat it.
Choice B: Climate change is a natural occurrence and policies should focus on economic growth.

Base Embedder: Everybody is responsible for dealing with climate change.

Llama 8B: A

Example of a fine-tuned embedder | Figures QA

How should society respond to climate change?
Choice A: Human activities are primarily responsible for climate change and urgent action is required to combat it.
Choice B: Climate change is a natural occurrence and policies should focus on economic growth.

Fine-tuned Embedder: Requiring corporations to disclose their climate impacts would be an unnecessary and counterproductive burden. Companies are already highly incentivized to reduce their environmental footprint in order to cut costs, [...]

Llama 8B: B

Gemma 9B: A